

Mridul Gulati

AI Technical Lead

Delhi NCR, India · [linkedin.com/in/mridul-gulati](https://www.linkedin.com/in/mridul-gulati) · github.com/Mridul-Gulati

SUMMARY

AI Technical Lead specialising in multi-agent orchestration, agent runtime security and LLM evaluation. I design end-to-end agent pipelines with LangGraph and Google ADK, and have scaled them to 30,000+ daily runs while cutting latency by over 50%. I ship agents with guardrails, token-cost observability and automated quality benchmarks in place before release.

EXPERIENCE

AI Technical Lead · HCLTech

Mar 2026 – Present · Noida, India

- Built an intelligent SRE agent integrating ServiceNow and Datadog (via MCP) for automated six-step root cause analysis and remediation planning, taking it from requirements to client demo in under three weeks.
- Cut the SRE agent's data-gathering phase by 55% with tool caching; hardened input/output guardrails with callbacks and Model Armor.
- Implemented pre-production agent evaluation and simulation using custom quality metrics and simulated user personas.
- Added per-agent token cost observability, optimising expensive turns and reporting through BigQuery.
- Designed a multi-agent candidate screening system: CV and JD parsing, match scoring, question generation and a first-round AI interview agent on Google Meet.
- Built and deployed 50+ enterprise-grade agents to production, with guardrails and observability in place before release.
- Worked with presales to turn new agent capabilities into client-facing solutions; interviewed mid-to-senior GenAI engineering candidates.

Award: Trailblazer Award, Q2 2026: highest quarterly recognition from the CTO and Ecosystems team, across 500+ employees.

Data Scientist · Turing

Sep 2024 – Jan 2026 · Remote

- Built an agent that reviews LLM training notebooks and scores hallucination, accuracy, instruction following and language quality, returning a score, a done/rework decision and specific feedback.
- Grounded the agent with RAG on client-approved reference notebooks to reduce hallucination; scaled it to 30,000+ runs per day.
- Adopted voluntarily by data trainers as a pre-submission check, cutting average review time from ~60 to 30–40 minutes, with 98% satisfaction across 1,000 reviewers.
- Re-architected the system from LangGraph to Google ADK after ADK's release.
- Promoted to Pod Lead within six months; led 5–6 AI engineers, owned code review and onboarding.
- Worked with client stakeholders to define the evaluation criteria the agent enforced.

EDUCATION

Guru Gobind Singh Indraprastha University

2021 – 2025 · New Delhi

B.Tech in Computer Science and Engineering, minor in AI/ML · CGPA 8.99

SKILLS

Agents and orchestration: LangGraph, Google ADK, LangChain, MCP, A2A, A2UI

Evaluation and safety: LLM-as-judge evals, Simulated personas, Guardrails, Model Armor, LangSmith

Models and platforms: Gemini, Vertex AI, Agent Runtime, Gemini Enterprise, Memory Bank, Document AI

Engineering: Python, SQL, FastAPI, Docker, Cloud Run, BigQuery, Firestore, Streamlit, Playwright

Integrations: ServiceNow, Datadog

CERTIFICATIONS

- Google Cloud Professional Machine Learning Engineer
- Google Professional Agentic AI Architect
- Gemini Enterprise Forward Deployed Engineer program
- Google Cloud Partner Specialist (CPS)
- 18 Google Cloud skill badges